

Edited_05 Amazon Fine Food Reviews Analysis_Logistic Regression

April 4, 2019

1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective: Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

2 [1]. Reading Data

2.1 [1.1] Loading the data

The dataset is available in two forms 1. .csv file 2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [198]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

In [199]: # using SQLite Table to read data.
con = sqlite3.connect('database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points
# you can change the number to any other number based on your computing power
```

```
# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 200000
# for tsne assignment you can take 5k data points
```

```
filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 200000
```

```
# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(0)
```

```
def partition(x):
    if x < 3:
        return 0
    return 1
```

```
#changing reviews with score less than 3 to be positive and vice-versa
```

```
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)
```

Number of data points in our data (200000, 10)

```
Out[199]:
```

	Id	ProductId	UserId	ProfileName	\
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	
2	3	B000LQOCHO	ABXLMWJIXXAIN	Natalia Corres	"Natalia Corres"

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	1	1	1	1303862400	
1	0	0	0	1346976000	
2	1	1	1	1219017600	

	Summary	Text
0	Good Quality Dog Food	I have bought several of the Vitality canned d...
1	Not as Advertised	Product arrived labeled as Jumbo Salted Peanut...
2	"Delight" says it all	This is a confection that has been around a fe...

```
In [200]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)
```

```
In [201]: print(display.shape)
display.head()
```

(80668, 7)

```
Out [201]:
```

	UserId	ProductId	ProfileName	Time	Score	\
0	#oc-R115TNMSPFT9I7	B005ZBZLT4	Breyton	1331510400	2	
1	#oc-R11D9D7SHXIJB9	B005HG9ESG	Louis E. Emory "hoppy"	1342396800	5	
2	#oc-R11DNU2NBKQ23Z	B005ZBZLT4	Kim Cieszykowski	1348531200	1	
3	#oc-R1105J5ZVQE25C	B005HG9ESG	Penguin Chick	1346889600	5	
4	#oc-R12KPBODL2B5ZD	B0070SBEV0	Christopher P. Presta	1348617600	1	

	Text	COUNT(*)
0	Overall its just OK when considering the price...	2
1	My wife has recurring extreme muscle spasms, u...	3
2	This coffee is horrible and unfortunately not ...	2
3	This will be the bottle that you grab from the...	3
4	I didnt like this coffee. Instead of telling y...	2

```
In [202]: display[display['UserId']=='AZY10LLTJ71NX']
```

```
Out [202]:
```

	UserId	ProductId	ProfileName	Time	\
80638	AZY10LLTJ71NX	B001ATMQK2	undertheshrine "undertheshrine"	1296691200	

	Score	Text	COUNT(*)
80638	5	I bought this 6 pack because for the price tha...	5

```
In [203]: display['COUNT(*)'].sum()
```

```
Out [203]: 393063
```

3 [2] Exploratory Data Analysis

3.1 [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [204]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

```
Out [204]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	\
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	

	HelpfulnessDenominator	Score	Time \
0	2	5	1199577600
1	2	5	1199577600
2	2	5	1199577600
3	2	5	1199577600
4	2	5	1199577600

	Summary \
0	LOACKER QUADRATINI VANILLA WAFERS
1	LOACKER QUADRATINI VANILLA WAFERS
2	LOACKER QUADRATINI VANILLA WAFERS
3	LOACKER QUADRATINI VANILLA WAFERS
4	LOACKER QUADRATINI VANILLA WAFERS

	Text
0	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
3	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
4	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8) ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [205]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)

In [206]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first')
final.shape

Out[206]: (160178, 10)

In [207]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100

Out[207]: 80.089
```

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [208]: display= pd.read_sql_query("""
```

```
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)
```

```
display.head()
```

```
Out[208]:
```

	Id	ProductId	UserId	ProfileName	\
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens	"Jeanne"
1	44737	B001EQ55RW	A2V0I904FH7ABY		Ram

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	3	1	5	1224892800	
1	3	2	4	1212883200	

	Summary	\
0	Bought This for My Son at College	
1	Pure cocoa taste with crunchy almonds inside	

	Text
0	My son loves spaghetti so I didn't hesitate or...
1	It was almost a 'love at first bite' - the per...

```
In [209]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [210]: #Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)
```

```
#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()
```

```
(160176, 10)
```

```
Out[210]:
```

1	134799
0	25377

Name: Score, dtype: int64

```
In [211]: ##Sorting data for Time Based Splitting
time_sorted_data = final.sort_values('Time', axis=0, ascending=True, inplace=False,

final = time_sorted_data.take(np.random.permutation(len(final))[:100000])
print(final.shape)
final.head()
```

(100000, 10)

```
Out[211]:
```

	Id	ProductId	UserId	ProfileName	\
	10291	11237	B001KVPCOG	AY74M03WTAOMB	Nut Nut
	79814	86781	B002DHBT7Q	A1CHKAWX7FAOM4	L. Kirk "Crabseye"
	76651	83391	B005ZBZLT4	A2QOYXPT6POXQS	Tony Barnes
	40580	44095	B00168ACG2	A026QTL5I5JRF	Suzanne Davis
	64314	69846	B002B8ODPW	A2ZE8BSZ5MMEOP	Jasmine "Uniquely Yours"

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
	10291	0	0	1	1226534400
	79814	2	2	1	1317600000
	76651	0	0	1	1349049600
	40580	0	0	1	1346112000
	64314	0	0	1	1343520000

	Summary	\
	10291	Best Roasted Almonds
	79814	One of my favorite cookies
	76651	Great deal!
	40580	THIS IS THE BEST CORNBREAD!!!!
	64314	Convenient and tasty

	Text
	10291 I've tried several other brands of roasted sal...
	79814 A cookie a day (or 2, depending on serving siz...
	76651 This coffee is a great deal! All the coffee d...
	40580 I have been eating this cornbread for years an...
	64314 Not my favorite, but for the price great produ...

4 [3] Preprocessing

4.1 [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [212]: # printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

```
I've tried several other brands of roasted salted almonds available via the Internet and in stores.
=====
The stigma of decafs, in general, has vanished! Emeril's Jazzed Up Decaf is, by far, the richest.
=====
This is not a traditional cookie! However, it is a good thing: it is its own little niche of yummies.
=====
Bought these for my grandbabies and they love them. I love them because they are organic. Only
=====
```

```
In [213]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
sent_1500 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)
```

```
I've tried several other brands of roasted salted almonds available via the Internet and in stores.
```

```
In [214]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all-urls
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
```



```

print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)

```

I've tried several other brands of roasted salted almonds available via the Internet and in stores. =====
The stigma of decafs, in general, has vanished! Emeril's Jazzed Up Decaf is, by far, the richest. =====
This is not a traditional cookie! However, it is a good thing: it is its own little niche of yours. =====
Bought these for my grandbabies and they love them. I love them because they are organic. Only

In [215]: # <https://stackoverflow.com/a/47091490/4084039>
import re

```

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase

```

In [216]: sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("="*50)

This is not a traditional cookie! However, it is a good thing: it is its own little niche of yours. =====

In [217]: #remove words with numbers python: <https://stackoverflow.com/a/18082370/4084039>
sent_0 = re.sub(r"\S*\d\S*", "", sent_0).strip()
print(sent_0)

I've tried several other brands of roasted salted almonds available via the Internet and in stores.

```
In [218]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)
```

This is not a traditional cookie. However, it is a good thing; it is its own little niche of yumminess.

```
In [219]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have been removed in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves',
                'you'll', 'you'd', 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
                'she', 'she's', 'her', 'hers', 'herself', 'it', 'it's', 'its', 'itself',
                'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that',
                'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has',
                'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as',
                'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through',
                'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off',
                'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'each',
                'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too',
                's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've",
                've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",
                "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'm',
                "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't",
                'won', "won't", 'wouldn', "wouldn't"])
```

```
In [220]: # Combining all the above students
from tqdm import tqdm
prepr_rev = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
    prepr_rev.append(sentence.strip())
```

100%| 100000/100000 [00:46<00:00, 2145.59it/s]

```
In [221]: print(len(prepr_rev))
          final.shape
```

```
100000
```

```
Out[221]: (100000, 10)
```

```
In [222]: prepr_rev[1500]
```

```
Out[222]: 'not traditional cookie however good thing little niche yumminess fig flavor subtle v
```

```
In [223]: final ['prepr_rev']= prepr_rev
          final.head(5)
```

```
Out[223]:
```

	Id	ProductId	UserId	ProfileName	\
10291	11237	B001KVPCOG	AY74M03WTAOMB	Nut Nut	
79814	86781	B002DHBT7Q	A1CHKAWX7FAOM4	L. Kirk "Crabseye"	
76651	83391	B005ZBZLT4	A2Q0YXPT6POXQS	Tony Barnes	
40580	44095	B00168ACG2	A026QTL5I5JRF	Suzanne Davis	
64314	69846	B002B8ODPW	A2ZE8BSZ5MMEOP	Jasmine "Uniquely Yours"	

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
10291	0	0	1	1226534400	
79814	2	2	1	1317600000	
76651	0	0	1	1349049600	
40580	0	0	1	1346112000	
64314	0	0	1	1343520000	

	Summary	\
10291	Best Roasted Almonds	
79814	One of my favorite cookies	
76651	Great deal!	
40580	THIS IS THE BEST CORNBREAD!!!!	
64314	Convenient and tasty	

	Text	\
10291	I've tried several other brands of roasted sal...	
79814	A cookie a day (or 2, depending on serving siz...	
76651	This coffee is a great deal! All the coffee d...	
40580	I have been eating this cornbread for years an...	
64314	Not my favorite, but for the price great produ...	

	prepr_rev
10291	tried several brands roasted salted almonds av...
79814	cookie day depending serving size one indulgen...
76651	coffee great deal coffee drinkers house enjoy ...
40580	eating cornbread years could not find locally ...
64314	not favorite price great product handy stores ...

[3.2] Preprocessing Review Summary

In [224]: *## Similarly you can do preprocessing for review summary also.*

```
preprocessed_summary = []
# tqdm is for printing the status bar
for summary in tqdm(final['Summary'].values):
    summary = re.sub(r"http\S+", "", summary) # remove urls from text python: https://
    summary = BeautifulSoup(summary, 'lxml').get_text() # https://stackoverflow.com/
    summary = decontracted(summary)
    summary = re.sub("\S*\d\S*", "", summary).strip() #remove words with numbers pyt
    summary = re.sub('[^A-Za-z]+', ' ', summary) #remove spacial character: https://
    # https://gist.github.com/sebleier/554280
    summary = ' '.join(e.lower() for e in summary.split() if e.lower() not in stopwords)
    preprocessed_summary.append(summary.strip())
```

100%|| 100000/100000 [00:29<00:00, 3428.23it/s]

```
In [225]: prepr_rev = [i + ' ' + j for i, j in zip(prepr_rev,preprocessed_summary)]
          print(prepr_rev[1500])
```

not traditional cookie however good thing little niche yumminess fig flavor subtle want full f

5 [4] Featurization

5.1 [4.1] BAG OF WORDS

```
In [226]: X = np.array(prepr_rev)
          y = np.array(final['Score'])
```

```
In [227]: from sklearn.model_selection import train_test_split
          #splitting data into Train, C.V and Test
          X_train, X_test, y_train, y_test = train_test_split(final ['prepr_rev'], final['Score'],
          X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test_size=0.33)
          print("Train:",X_train.shape,y_train.shape)
          print("CV:",X_cv.shape,y_cv.shape)
          print("Test:",X_test.shape,y_test.shape)
```

Train: (44890,) (44890,)

CV: (22110,) (22110,)

Test: (33000,) (33000,)

```
In [228]: vectorizer = CountVectorizer(min_df=10, max_features=500)
          vectorizer.fit(X_train)
          #vectorizer.fit(X_train) # fit has to happen only on train data
          # we use the fitted CountVectorizer to convert the text to vector
```

```

X_train_bow = vectorizer.transform(X_train)
X_cv_bow = vectorizer.transform(X_cv)
X_test_bow = vectorizer.transform(X_test)
print("After vectorizations")
print(X_train_bow.shape, y_train.shape)
print(X_cv_bow.shape, y_cv.shape)
print(X_test_bow.shape, y_test.shape)

```

```

After vectorizations
(44890, 500) (44890,)
(22110, 500) (22110,)
(33000, 500) (33000,)

```

5.2 [4.3] TF-IDF

```

In [229]: tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10, max_features=500)
          tf_idf_vect.fit(X_train)
          # we use the fitted CountVectorizer to convert the text to vector
          X_train_tfidf = tf_idf_vect.transform(X_train)
          X_cv_tfidf = tf_idf_vect.transform(X_cv)
          X_test_tfidf = tf_idf_vect.transform(X_test)
          print("After vectorizations")
          print(X_train_tfidf.shape, y_train.shape)
          print(X_cv_tfidf.shape, y_cv.shape)
          print(X_test_tfidf.shape, y_test.shape)

```

```

After vectorizations
(44890, 500) (44890,)
(22110, 500) (22110,)
(33000, 500) (33000,)

```

5.3 [4.4] Word2Vec

```

In [230]: # List of sentence in X_train text
          sent_of_train=[]
          for sent in X_train:
              sent_of_train.append(sent.split())

          # List of sentence in X_test text
          sent_of_test=[]
          for sent in X_test:
              sent_of_test.append(sent.split())

          # Train your own Word2Vec model using your own train text corpus
          # min_count = 5 considers only words that occurred at least 5 times
          w2v_model=Word2Vec(sent_of_train,min_count=5,size=50, workers=4)

```

```

print(w2v_model.wv.most_similar('great'))
print('='*50)
print(w2v_model.wv.most_similar('worst'))

w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_words))

[('fantastic', 0.8478500843048096), ('good', 0.8255733251571655), ('awesome', 0.81360322237014)
=====
[('best', 0.7517200708389282), ('nastiest', 0.7383694648742676), ('greatest', 0.71207940578460)
number of words that occurred minimum 5 times 12960

```

```

In [231]: w2v_words = list(w2v_model.wv.vocab)
          print("number of words that occurred minimum 5 times ",len(w2v_words))
          print("sample words ", w2v_words[0:50])

number of words that occurred minimum 5 times 12960
sample words ['tea', 'go', 'everyday', 'nearly', 'indestructible', 'forgotten', 'steeped', 'm

```

5.4 [4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

[4.4.1.1] Avg W2v

```

In [166]: i=0
          sent_of_test_cv=[]
          for sentence in X_cv:
              sent_of_test_cv.append(sentence.split())

In [167]: # average Word2Vec
          # compute average word2vec for each review.
          sent_vectors_cv = []; # the avg-w2v for each sentence/review is stored in this list
          for sent in tqdm(sent_of_test_cv): # for each review/sentence
              sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need
              cnt_words = 0; # num of words with a valid vector in the sentence/review
              for word in sent: # for each word in a review/sentence
                  if word in w2v_words:
                      vec = w2v_model.wv[word]
                      sent_vec += vec
                      cnt_words += 1
              if cnt_words != 0:
                  sent_vec /= cnt_words
              sent_vectors_cv.append(sent_vec)
          sent_vectors_cv = np.array(sent_vectors_cv)
          print(sent_vectors_cv.shape)
          print(sent_vectors_cv[0])

```

100%|| 22110/22110 [01:05<00:00, 336.99it/s]

(22110, 50)

```
[-0.52553881 -0.56571381  0.63142397  0.43928345  0.59151121  0.34461575
 0.18870962  0.45058656 -0.82038542 -0.29991521  0.56942672 -0.67313881
-0.71965228 -0.07363756  0.2739284   0.0503723   0.52536564  0.63691369
 0.18458937  0.23666228 -0.32867154 -0.15568994  0.09874957 -0.01075351
 0.33511867  1.14519034  0.29567677 -0.0468841  -0.03616501  0.08883335
-0.36886215 -0.42454403 -0.57793536  0.37274055 -0.06555429  0.18378313
 0.30609365 -0.04986344  0.03503402  0.93081964 -0.07000684 -0.54992588
 0.2967225   0.18228281  0.0962322   0.48645934 -0.23933173  0.25615076
 0.29496165 -0.3852263 ]
```

In [168]: *# compute average word2vec for X_test .*

```
test_vectors = [];
for sent in tqdm(sent_of_test):
    sent_vec = np.zeros(50)
    cnt_words =0;
    for word in sent: #
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    test_vectors.append(sent_vec)

test_vectors = np.array(test_vectors)

print(test_vectors.shape)
print(test_vectors[0])
```

100%|| 33000/33000 [01:37<00:00, 338.26it/s]

(33000, 50)

```
[-0.19294365 -0.66535483  0.49491829  0.61472203  0.43786552 -0.2212762
 0.63726501 -0.09604544 -1.01787599 -0.38894434 -0.02899949 -0.36255753
-0.46169367 -0.24737216 -0.55879127  0.22363069  0.50083397  0.24539507
-0.0535983   0.33607479  0.03222066 -0.15921549  0.32272275  0.31025571
 0.21540802  0.19478748 -0.61918377 -0.0017801  -0.71217188 -0.46075358
-0.54823264 -0.17418158 -0.12233295 -0.03316205 -0.08219504  0.14051325
-0.26644893  0.27429939 -0.4835991   0.73440272  0.55616765 -0.06391602
-0.40270575 -0.12931861 -0.24423024  0.236344   0.26991326  0.16444251
-0.22399055  0.54960559]
```

```
In [169]: # compute average word2vec for X_train .
```

```
train_vectors = [];  
for sent in tqdm(sent_of_train):  
    sent_vec = np.zeros(50)  
    cnt_words = 0;  
    for word in sent: #  
        if word in w2v_words:  
            vec = w2v_model.wv[word]  
            sent_vec += vec  
            cnt_words += 1  
    if cnt_words != 0:  
        sent_vec /= cnt_words  
    train_vectors.append(sent_vec)  
  
train_vectors = np.array(train_vectors)  
  
print(train_vectors.shape)  
print(train_vectors[0])
```

```
100%|| 44890/44890 [02:09<00:00, 346.49it/s]
```

```
(44890, 50)
```

```
[-0.14060537 -0.72614364  0.32810215  1.11755044  0.23384139  0.28288134  
-0.01866034 -0.04290673 -0.52098112 -0.01752026 -0.01945195  0.15767466  
-0.07230324 -0.01717835 -0.27854393  0.05381443  0.73443291  0.75394678  
 0.93170925  0.04095242 -0.49682882 -0.20451759  0.2238511  0.65370598  
-0.13699235 -0.40837955  0.57271826 -0.12713724 -0.59230489 -0.02618843  
-0.34113196 -0.33512895 -0.13026007 -0.22999276  0.30137488  0.37416903  
 0.2087824  0.68820576  0.2217927  0.66206925  0.37959712  0.22534903  
-0.7366674 -0.62474787  0.04708067  0.12064839  0.31486852  0.38119362  
 0.3194624 -0.3280405 ]
```

[4.4.1.2] TFIDF weighted W2v

```
In [170]: tf_idf_vect = TfidfVectorizer()
```

```
# final_tf_idf1 is the sparse matrix with row= sentence, col=word and cell_val = tfi  
final_tf_idf1 = tf_idf_vect.fit_transform(X_train)  
dictionary = dict(zip(tf_idf_vect.get_feature_names(), list(tf_idf_vect.idf_)))  
  
# tfidf words/col-names  
tfidf_feat = tf_idf_vect.get_feature_names()  
  
# compute TFIDF Weighted Word2Vec for X_test .  
tfidf_test_vectors = [];
```



```

row=0;
for sent in tqdm(sent_of_test):
    sent_vec = np.zeros(50)
    weight_sum =0;
    for word in sent:
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_test_vectors.append(sent_vec)

tfidf_test_vectors = np.array(tfidf_test_vectors)
print(tfidf_test_vectors.shape)
print(tfidf_test_vectors[0])

```

100%|| 33000/33000 [22:17<00:00, 24.67it/s]

(33000, 50)

```

[-0.16686663 -0.86323916  0.49690235  0.5184359   0.56615715 -0.54148702
  0.93596545 -0.25554939 -1.12918029 -0.77515194  0.04043236 -0.01329734
 -0.5751481  -0.6871599  -0.75427953  0.1534776   0.36420995  0.26531452
  0.08081806  0.42276238  0.15202629 -0.0840503   0.46600952  0.41430244
  0.26516221 -0.03617444 -0.87445306 -0.01075618 -0.63894613 -0.59917669
 -0.83022163 -0.19207686  0.0749336  -0.06149395 -0.1844047   0.28045885
 -0.81266072  0.22030619 -0.57452736  0.7226248   0.6758246   0.15453529
 -0.08191958  0.15770662 -0.32205749  0.08399961  0.39292019  0.16717439
 -0.39050156  0.80903407]

```

```

In [171]: # compute TFIDF Weighted Word2Vec for X_train .
tfidf_train_vectors = [];
row=0;
for sent in tqdm(sent_of_train):
    sent_vec = np.zeros(50)
    weight_sum =0;
    for word in sent:
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_train_vectors.append(sent_vec)

```

```
tfidf_train_vectors = np.array(tfidf_train_vectors)
print(tfidf_train_vectors.shape)
print(tfidf_train_vectors[0])
```

100%|| 44890/44890 [27:49<00:00, 26.89it/s]

(44890, 50)

```
[-0.29540099 -0.85954222 -0.18179773  0.88866737  0.38733377  0.07670469
 -0.08424611  0.12395096 -0.48383716 -0.07993022 -0.00500073  0.31535502
  0.06399823  0.02668346 -0.39874532 -0.19141424  0.41162859  0.67405128
  1.09345726  0.07239559 -0.44772274  0.12750844  0.26261354  0.38511071
 -0.16009767 -0.39785692  1.16591441 -0.4437846  -0.38423109  0.32179244
 -0.11075185 -0.21630509  0.073712  -0.27370394  0.21916168  0.2249126
 -0.02694394  0.95540116  0.0286596  0.38618412  0.51026674  0.42937773
 -0.9517379  -0.66895635  0.19392165 -0.37172587  0.3754621  0.14247962
  0.01241208 -0.48043844]
```

6 [5] Assignment 5: Apply Logistic Regression

- Apply Logistic Regression on these feature sets**

-

- SET 1:**Review text, preprocessed one converted into vectors

- SET 2:**Review text, preprocessed one converted into vectors

- SET 3:**Review text, preprocessed one converted into vectors

- SET 4:**Review text, preprocessed one converted into vectors

-

- Hyper paramter tuning (find best hyper parameters corresponding the algorithm that**

-

- Find the best hyper parameter which will give the maximum <https://www.appliedaicom>

- Find the best hyper paramter using k-fold cross validation or simple cross validation data

- Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this task

-

- Pertubation Test**

-

- Get the weights W after fit your model with the data X i.e Train data.

- Add a noise to the X ($X' = X + e$) and get the new data set X' (if X is a sparse

matrix, $X.data += e$)

- Fit the model again on data X' and get the weights W'

- Add a small eps value(to eliminate the divisibile by zero error) to W and W i.e

$W = W + 10^{-6}$ and $W' = W' + 10^{-6}$

```
<li>Now find the % change between W and W' (| (W-W') / (W) |)*100</li>
<li>Calculate the 0th, 10th, 20th, 30th, ...100th percentiles, and observe any sudden rise in
<li> Ex: consider your 99th percentile is 1.3 and your 100th percentiles are 34.6, there is su
    <li> Print the feature names whose % change is more than a threshold x(in our example :
    </ul>
</li>
<br>
<li><strong>Sparsity</strong>
    <ul>
<li>Calculate sparsity on weight vector obtained after using L1 regularization</li>
    </ul>
</li>
<br><font color='red'>NOTE: Do sparsity and multicollinearity for any one of the vectorizers.
<br>
<br>
<li><strong>Feature importance</strong>
    <ul>
<li>Get top 10 important features for both positive and negative classes separately.</li>
    </ul>
</li>
<br>
<li><strong>Feature engineering</strong>
    <ul>
<li>To increase the performance of your model, you can also experiment with with feature engin
        <ul>
<li>Taking length of reviews as another feature.</li>
<li>Considering some features from review summary as well.</li>
        </ul>
    </ul>
</li>
<br>
<li><strong>Representation of results</strong>
    <ul>
<li>You need to plot the performance of model both on train data and cross validation data for
<img src='train_cv_auc.JPG' width=300px></li>
<li>Once after you found the best hyper parameter, you need to train your model with it, and f
<img src='train_test_auc.JPG' width=300px></li>
<li>Along with plotting ROC curve, you need to print the <a href='https://www.appliedaicourse.
<img src='confusion_matrix.png' width=300px></li>
    </ul>
</li>
<br>
<li><strong>Conclusion</strong>
    <ul>
<li>You need to summarize the results at the end of the notebook, summarize it in the table for
    <img src='summary.JPG' width=400px>
```


Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakag, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method fit_transform() on you train data, and apply the method transform() on cv/test data.
4. For more details please go through this link.

7 Applying Logistic Regression

7.1 [5.1] Logistic Regression on BOW, SET 1

7.1.1 [5.1.1] Applying Logistic Regression with L1 regularization on BOW, SET 1

[5.1.1.1] Calculating sparsity on weight vector obtained using L1 regularization on BOW, SET 1

```
In [232]: # Please write all the code with proper documentation
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_auc_score
```

```
def lr_all (X_train,y_train,X_cv,penal):
    train_auc = []
    cv_auc = []
    hyper_param = [0.001, 0.01, 0.1, 1, 10, 100,1000]
```

```
    for i in tqdm(hyper_param):
        lr = LogisticRegression(C=i,penalty= penal)
        lr.fit(X_train_bow,y_train)
```

```
    y_train_pred = lr.predict_proba(X_train_bow)[:,-1]
    y_cv_pred = lr.predict_proba(X_cv_bow)[:,-1]
```

```
    train_auc.append(roc_auc_score(y_train,y_train_pred))
    cv_auc.append(roc_auc_score(y_cv, y_cv_pred))
```

```
#Plotting the AUC curve for train and test data
```

```
#Print the graph
```

```
# https://stackoverflow.com/questions/28077499/matplotlib-pyplot-plot-x-axis-tick-labels
```

```
plt.plot(np.log(hyper_param), train_auc, label='Train AUC-L1')
plt.plot(np.log(hyper_param), cv_auc, label='CV AUC-L1')
```

```

plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left', borderaxespad=0.)
plt.grid(True)
plt.legend()
plt.xlabel("Lambda: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
#Cv-auc scores with penalty L1
print("-----")
print("Cv auc scores with penalty L1")
print(cv_auc)
print("Maximun Auc value :",max(cv_auc))
print("Index",cv_auc.index(max(cv_auc)))

#Get lambda value for max auc in cv data
mx = 0
for i in range(len(cv_auc)):
    if(cv_auc[i]> cv_auc[mx]):
        mx = i
best = hyper_param[mx]
print("The optimal value of Lambda = ", best)

lr = LogisticRegression(C =best,penalty= 'l1', class_weight = 'balanced')
lr.fit(X_train_bow,y_train)

train_fpr, train_tpr, thresholds = roc_curve(y_train, lr.predict_proba(X_train_bow)[0,:])
test_fpr, test_tpr, thresholds = roc_curve(y_test, lr.predict_proba(X_test_bow)[0,:])

plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))

plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left', borderaxespad=0.)
plt.grid(True)
plt.legend()
plt.xlabel("Lamda: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

#Confusion Matrix

print("Train confusion matrix")
print(confusion_matrix(y_train, lr.predict(X_train_bow)))

```

```

print("Test confusion matrix")
print(confusion_matrix(y_test, lr.predict(X_test_bow)))

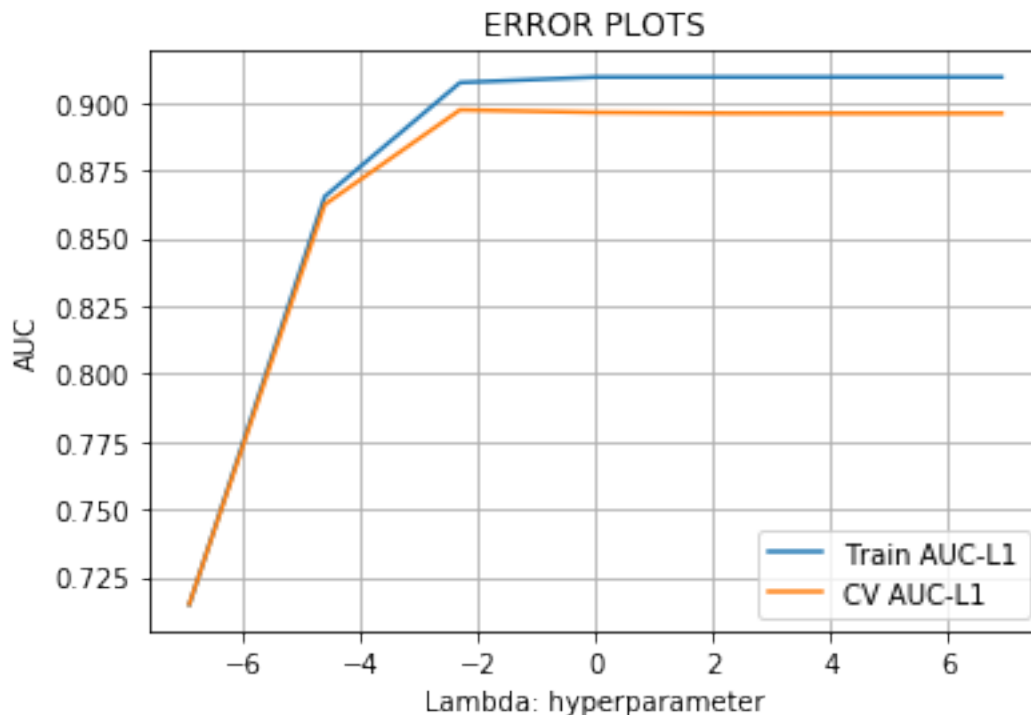
cm = confusion_matrix(y_train, lr.predict(X_train_bow))
cm = confusion_matrix(y_test, lr.predict(X_test_bow))
tn, fp, fn, tp = cm.ravel()
# https://stackoverflow.com/questions/35572000/how-can-i-plot-a-confusion-matrix
# Code for drawing seaborn heatmaps
class_names = ['0', '1']
df_heatmap = pd.DataFrame(cm, index=class_names, columns=class_names )
fig = plt.figure(figsize=(5,3))
heatmap = sns.heatmap(df_heatmap, annot=True, fmt="d")

# Setting tick labels for heatmap
heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0, ha='right')
heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=0, ha='right')
plt.ylabel('True label',size=18)
plt.xlabel('Predict label',size=18)
plt.title("Confusion Matrix\n",size=24)
plt.show()

```

In [233]: lr_all(X_train_bow,y_train,X_cv_bow,'l1')

100%|| 7/7 [00:02<00:00, 2.96it/s]




Cv auc scores with penalty L1

[0.7149699199648194, 0.8623282056497347, 0.8972412121357145, 0.8963105754704795, 0.89600729675]

Maximun Auc value : 0.8972412121357145

Index 2

The optimal value of Lambda = 0.1



Train confusion matrix

```
[[ 6174 1081]
```


```
 [ 6539 31096]]
```

Test confusion matrix

```
[[ 4191   917]
```

```
 [ 4831 23061]]
```

Confusion Matrix



7.1.2 [5.1.2] Applying Logistic Regression with L2 regularization on BOW, SET 1

```
In [236]: def all_lr (X_train,y_train,X_cv,penal):
    train_auc_l2 = []
    cv_auc_l2 = []
    hyper_param = [0.001, 0.01, 0.1, 1, 10, 100,1000]

    for i in tqdm(hyper_param):
        lr = LogisticRegression(C=i,penalty= penal)
        lr.fit(X_train_bow,y_train)

        y_train_pred = lr.predict_proba(X_train_bow)[:,-1]
        y_cv_pred = lr.predict_proba(X_cv_bow)[:,-1]

        train_auc_l2.append(roc_auc_score(y_train,y_train_pred))
        cv_auc_l2.append(roc_auc_score(y_cv, y_cv_pred))

    #Plotting the AUC curve for train and test data

    #Print the graph

    # https://stackoverflow.com/questions/28077499/matplotlib-pyplot-plot-x-axis-tick-labels
```



```

plt.plot(np.log(hyper_param), train_auc_l2, label='Train AUC-L2')
plt.plot(np.log(hyper_param), cv_auc_l2, label='CV AUC-L2')
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left', borderaxespad=0.)
plt.grid(True)
plt.legend()
plt.xlabel("Lambda: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

#Cv auc scores with penalty L2
print("-----")
print("Cv auc scores with penalty L2")
print(cv_auc_l2)
print("Maximun Auc value :",max(cv_auc_l2))
print("Index",cv_auc_l2.index(max(cv_auc_l2)))


#Get lambda value for max auc in cv data
mx = 0
for i in range(len(cv_auc_l2)):
    if(cv_auc_l2[i]> cv_auc_l2[mx]):
        mx = i
best = hyper_param[mx]
print("The optimal value of Lambda = ", best)


lr = LogisticRegression(C= best, penalty= 'l2', class_weight = 'balanced')
lr.fit(X_train_bow,y_train)

train_fpr, train_tpr, thresholds = roc_curve(y_train, lr.predict_proba(X_train_bow)[0])
test_fpr, test_tpr, thresholds = roc_curve(y_test, lr.predict_proba(X_test_bow)[0])

plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left', borderaxespad=0.)
plt.grid(True)
plt.legend()
plt.xlabel("Lamda: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

```

#Confusion Matrix

```
print("Train confusion matrix")
print(confusion_matrix(y_train, lr.predict(X_train_bow)))
print("Test confusion matrix")
print(confusion_matrix(y_test, lr.predict(X_test_bow)))
```

```
cm = confusion_matrix(y_train, lr.predict(X_train_bow))
```

```
cm = confusion_matrix(y_test, lr.predict(X_test_bow))
```

```
tn, fp, fn, tp = cm.ravel()
```

```
# https://stackoverflow.com/questions/35572000/how-can-i-plot-a-confusion-matrix
```

```
# Code for drawing seaborn heatmaps
```

```
class_names = ['0', '1']
```

```
df_heatmap = pd.DataFrame(cm, index=class_names, columns=class_names )
```

```
fig = plt.figure(figsize=(5,3))
```

```
heatmap = sns.heatmap(df_heatmap, annot=True, fmt="d")
```

```
# Setting tick labels for heatmap
```

```
heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0, ha='right')
```

```
heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=0, ha='right')
```

```
plt.ylabel('True label',size=18)
```


```
plt.xlabel('Predict label',size=18)
```

```
plt.title("Confusion Matrix\n",size=24)
```

```
plt.show()
```

```
In [237]: all_lr(X_train_bow,y_train,X_cv_bow,'l2')
```

```
100%|| 7/7 [00:04<00:00, 1.47it/s]
```




Cv auc scores with penalty L2

[0.8686668268851596, 0.8946634880571093, 0.8969706221466804, 0.8961376090875506, 0.89598988413]

Maximun Auc value : 0.8969706221466804

Index 2

The optimal value of Lambda = 0.1




Train confusion matrix

```
[[ 6162 1093]
 [ 6475 31160]]
```

Test confusion matrix

```
[[ 4192  916]
 [ 4798 23094]]
```

Confusion Matrix



```
In [238]: from scipy.sparse import find
           #Weights before adding random noise
           weights1 = find(lr.coef_[0])[2]
           print(weights1[:10])

[ 0.15021456 -0.03929764 -0.05173317  0.43961426  0.05990978 -0.27028027
 -0.5902203  -0.13474628 -0.55136027  0.12501951]
```

[5.1.2.1] Performing pertubation test (multicollinearity check) on BOW, SET 1

```
In [239]: # Please write all the code with proper documentation

lr = LogisticRegression(C=10,penalty='l2')
lr.fit(X_train_bow,y_train)
weight1 = lr.coef_ # the weight vector

# New dataset by adding a small noise
new_train = X_train_bow.astype(float)
new_train.data += np.random.uniform(-0.0001,0.0001,1 )

# Fit the model on new data
lr = LogisticRegression(C=1,penalty='l2')
lr.fit(X_train_bow,y_train)
```

```

weight2 =lr.coef_

weight1 += 10**-6
weight2 += 10**-6

percentage_change_vector = abs( (weight1-weight2) / (weight1) )*100

```

```

In [240]: #print weights difference
print(percentage_change_vector.max())
print(percentage_change_vector.min())
print(percentage_change_vector.std())

```

```

82.51198275414353
0.013428872363584931
6.011225674585524

```

```

In [241]: percentage_change=[]
collinear_features=[]

for i in range(1,101):
    f=np.where(percentage_change_vector > i)[1].size
    percentage_change.append(i)
    collinear_features.append(f)

```

```

In [242]: feat = vectorizer.get_feature_names()
print("No of features have weight changes greater than 30%: ", percentage_change_vector)
fe=[]
print("\ncollinear features are :")
for i in np.where(percentage_change_vector > 1)[1]:
    fe.append(feat[i])
print(fe)

```

```

No of features have weight changes greater than 30%:  5

```

```

collinear features are :
['able', 'actually', 'added', 'aftertaste', 'ago', 'almost', 'already', 'alternative', 'although']

```

```

In [243]: t = range(0,101,10)
for i in t:
    print(i, "th percentile : ",np.percentile(percentage_change_vector,i))

plt.plot(t,np.percentile(percentage_change_vector,t) )

```

```

0 th percentile :  0.013428872363584931
10 th percentile :  0.3774116085177026


```

```

20 th percentile : 0.697804242923374
30 th percentile : 0.9233637864705805
40 th percentile : 1.1362068754828873
50 th percentile : 1.3431337469014863
60 th percentile : 1.5986589368506539
70 th percentile : 1.9457369935720714
80 th percentile : 2.3428249702073187
90 th percentile : 3.313082637600024
100 th percentile : 82.51198275414353

```

Out[243]: [matplotlib.lines.Line2D at 0x1a4742f0f0>]



```

In [244]: t = np.linspace(99,100,10)
          for i in t:
              print(i, "th percentile : ",np.percentile(percentage_change_vector,i))

          plt.plot(t,np.percentile(percentage_change_vector,t) )

```

```

99.0 th percentile : 19.63632149985406
99.11111111111111 th percentile : 31.68439678888494
99.22222222222223 th percentile : 41.535800817665084
99.33333333333333 th percentile : 43.339586771510895
99.44444444444444 th percentile : 45.06761464189339
99.55555555555556 th percentile : 46.68699433404016
99.66666666666667 th percentile : 54.62190417904557


```

```

99.77777777777777 th percentile : 66.64210745626136
99.88888888888889 th percentile : 75.02188514028092
100.0 th percentile : 82.51198275414353

```

Out[244]: [<matplotlib.lines.Line2D at 0x1a44599a20>]



7.1.3 [5.1.3] Feature Importance on BOW, SET 1

[5.1.3.1] Top 10 important features of positive class from SET 1

In [245]: *# Please write all the code with proper documentation*

```
all_features = vectorizer.get_feature_names()
```

```
def most_informative_features(vectorizer, lr, n= 25):
```

```
    feature_names = all_features
```

```
    coefs_with_fns = sorted(zip(lr.coef_[0], feature_names))
```

```
    top = zip(coefs_with_fns[:n], coefs_with_fns[n:])
```

```
    print("\t\tPositive\t\t\t\t\tNegative")
```

```
    print("-----")
```

```
    for (coef_1, fn_1), (coef_2, fn_2) in top:
```

```
        print("\t%.4f\t%-15s\t\t\t\t\t%.4f\t%-15s" % (coef_1, fn_1, coef_2, fn_2))
```

In [246]: most_informative_features(X_train_bow,lr)

Positive		Negative	

1.8518	amazing	-1.5299	disappoint
1.6625	perfect	-1.2900	money
1.6435	excellent	-1.0010	stick
1.5841	delicious	-0.7998	away
1.5318	pleased	-0.7797	thought
1.4619	highly	-0.7468	nothing
1.4241	smooth	-0.7252	guess
1.3956	wonderful	-0.6910	bad
1.2573	yummy	-0.6908	rather
1.2346	awesome	-0.6460	opened
1.2338	loves	-0.6136	reviews
1.2224	glad	-0.6096	tasted
1.2101	great	-0.5683	received
1.1767	surprised	-0.5632	left
1.1650	best	-0.5628	service
1.1544	thank	-0.5375	looked
1.0184	tasty	-0.5335	cannot
1.0117	exactly	-0.5204	maybe
1.0067	nice	-0.5130	instead
0.9683	stores	-0.5098	bought
0.9572	thanks	-0.5019	bitter
0.9096	definitely	-0.4868	gave
0.9057	happy	-0.4864	even
0.8527	easy	-0.4828	aftertaste
0.8428	works	-0.4777	ingredient

[5.1.3.2] Top 10 important features of negative class from SET 1

In [247]: *# Please write all the code with proper documentation*
most_informative_features(X_train_bow,lr)

Positive		Negative	

1.8518	amazing	-1.5299	disappoint
1.6625	perfect	-1.2900	money
1.6435	excellent	-1.0010	stick
1.5841	delicious	-0.7998	away
1.5318	pleased	-0.7797	thought
1.4619	highly	-0.7468	nothing
1.4241	smooth	-0.7252	guess
1.3956	wonderful	-0.6910	bad
1.2573	yummy	-0.6908	rather
1.2346	awesome	-0.6460	opened


1.2338	loves	-0.6136	reviews
1.2224	glad	-0.6096	tasted
1.2101	great	-0.5683	received
1.1767	surprised	-0.5632	left
1.1650	best	-0.5628	service
1.1544	thank	-0.5375	looked
1.0184	tasty	-0.5335	cannot
1.0117	exactly	-0.5204	maybe
1.0067	nice	-0.5130	instead
0.9683	stores	-0.5098	bought
0.9572	thanks	-0.5019	bitter
0.9096	definitely	-0.4868	gave
0.9057	happy	-0.4864	even
0.8527	easy	-0.4828	aftertaste
0.8428	works	-0.4777	ingredient

7.2 [5.2] Logistic Regression on TFIDF, SET 2

7.2.1 [5.2.1] Applying Logistic Regression with L1 regularization on TFIDF, SET 2

In [248]: `lr_all(X_train_tfidf,y_train,X_cv_tfidf,'l1')`

100%|| 7/7 [00:02<00:00, 3.18it/s]




Cv auc scores with penalty L1

[0.7149699199648194, 0.8623288923876233, 0.8972410900489787, 0.896310926469845, 0.896007037323]

Maximun Auc value : 0.8972410900489787

Index 2

The optimal value of Lambda = 0.1




Train confusion matrix

```
[[ 6174 1081]
 [ 6539 31096]]
```

Test confusion matrix

```
[[ 4191  917]
 [ 4831 23061]]
```


Confusion Matrix



7.2.2 [5.2.2] Applying Logistic Regression with L2 regularization on TFIDF, SET 2

```
In [249]: all_lr(X_train_tfidf,y_train,X_cv_tfidf,'l2')
```

```
100%|| 7/7 [00:04<00:00, 1.46it/s]
```




Cv auc scores with penalty L2

[0.8686668268851596, 0.8946634880571093, 0.8969706221466804, 0.8961376090875506, 0.89598988413

Maximun Auc value : 0.8969706221466804

Index 2

The optimal value of Lambda = 0.1




Train confusion matrix

```
[[ 6162 1093]
 [ 6475 31160]]
```

Test confusion matrix

```
[[ 4192  916]
 [ 4798 23094]]
```

Confusion Matrix



7.2.3 [5.2.3] Feature Importance on TFIDF, SET 2

[5.2.3.1] Top 10 important features of positive class from SET 2

In [250]: *# Please write all the code with proper documentation*

```
# Please write all the code with proper documentation
all_features = tf_idf_vect.get_feature_names()

def most_informative_features(tf_idf_vect, lr, n= 25):
    feature_names = all_features
    coefs_with_fns = sorted(zip(lr.coef_[0], feature_names))
    top = zip(coefs_with_fns[:n + 1]: -1], coefs_with_fns[:n])
    print("\t\tPositive\t\t\t\t\tNegative")
    print("-----")
    for (coef_1, fn_2), (coef_2, fn_1) in top:
        print("\t%.4f\t%-15s\t\t\t\t\t%.4f\t%-15s" % (coef_1, fn_2, coef_2, fn_1))
```

In [251]: `most_informative_features(X_train_tfidf,lr)`

	Positive		Negative	
	1.8518	amount	-1.5299	drinking

1.6625	pieces	-1.2900	never
1.6435	fan	-1.0010	stores
1.5841	dog	-0.7998	baby
1.5318	plus	-0.7797	three
1.4619	house	-0.7468	nuts
1.4241	smooth	-0.7252	happy
1.3956	work	-0.6910	bags
1.2573	yummy	-0.6908	read
1.2346	back	-0.6460	orange
1.2338	maybe	-0.6136	review
1.2224	got	-0.6096	tastes like
1.2101	gum	-0.5683	reason
1.1767	take	-0.5632	likes
1.1650	big	-0.5628	service
1.1544	thanks	-0.5375	makes
1.0184	tea	-0.5335	cats
1.0117	family	-0.5204	months
1.0067	not good	-0.5130	kind
0.9683	subscribe	-0.5098	boxes
0.9572	thing	-0.5019	bold
0.9096	disappointed	-0.4868	giving
0.9057	help	-0.4864	excellent
0.8527	end	-0.4828	ago
0.8428	worth	-0.4777	keurig

[5.2.3.2] Top 10 important features of negative class from SET 2

In [252]: *# Please write all the code with proper documentation*
`most_informative_features(X_train_tfidf,lr)`

Positive		Negative	

1.8518	amount	-1.5299	drinking
1.6625	pieces	-1.2900	never
1.6435	fan	-1.0010	stores
1.5841	dog	-0.7998	baby
1.5318	plus	-0.7797	three
1.4619	house	-0.7468	nuts
1.4241	smooth	-0.7252	happy
1.3956	work	-0.6910	bags
1.2573	yummy	-0.6908	read
1.2346	back	-0.6460	orange
1.2338	maybe	-0.6136	review
1.2224	got	-0.6096	tastes lik
1.2101	gum	-0.5683	reason
1.1767	take	-0.5632	likes

1.1650	big	-0.5628	service
1.1544	thanks	-0.5375	makes
1.0184	tea	-0.5335	cats
1.0117	family	-0.5204	months
1.0067	not good	-0.5130	kind
0.9683	subscribe	-0.5098	boxes
0.9572	thing	-0.5019	bold
0.9096	disappointed	-0.4868	giving
0.9057	help	-0.4864	excellent
0.8527	end	-0.4828	ago
0.8428	worth	-0.4777	keurig


7.3 [5.3] Logistic Regression on AVG W2V, SET 3

7.3.1 [5.3.1] Applying Logistic Regression with L1 regularization on AVG W2V SET 3

In [253]: *# Please write all the code with proper documentation*

```
lr_all(train_vectors,y_train,X_cv,'l1')
```

100%|| 7/7 [00:02<00:00, 3.00it/s]




Cv auc scores with penalty L1

[0.7149699199648194, 0.8623289381701493, 0.8972420820037071, 0.8963110790782647, 0.89600644215]

Maximun Auc value : 0.8972420820037071

Index 2

The optimal value of Lambda = 0.1



Train confusion matrix

[[6174 1081]


[6539 31096]]

Test confusion matrix

[[4191 917]

[4831 23061]]

Confusion Matrix



7.3.2 [5.3.2] Applying Logistic Regression with L2 regularization on AVG W2V, SET 3

```
In [254]: # Please write all the code with proper documentation
          all_lr(train_vectors,y_train,X_cv,'l2')
```

```
100%|| 7/7 [00:04<00:00, 1.43it/s]
```




Cv auc scores with penalty L2

[0.8686668268851596, 0.8946634880571093, 0.8969706221466804, 0.8961376090875506, 0.89598988413]

Maximun Auc value : 0.8969706221466804

Index 2

The optimal value of Lambda = 0.1




Train confusion matrix

```
[[ 6162 1093]
 [ 6475 31160]]
```

Test confusion matrix

```
[[ 4192  916]
 [ 4798 23094]]
```

Confusion Matrix




7.4 [5.4] Logistic Regression on TFIDF W2V, SET 4

7.4.1 [5.4.1] Applying Logistic Regression with L1 regularization on TFIDF W2V, SET 4

In [255]: *# Please write all the code with proper documentation*
lr_all(tfidf_train_vectors,y_train,X_cv,'l1')

100%|| 7/7 [00:02<00:00, 2.70it/s]




Cv auc scores with penalty L1

[0.7149699199648194, 0.8623300979941393, 0.897239838659937, 0.8963122846847806, 0.896006747367]

Maximun Auc value : 0.897239838659937

Index 2

The optimal value of Lambda = 0.1



Train confusion matrix

```
[[ 6174 1081]
```


```
 [ 6539 31096]]
```

Test confusion matrix

```
[[ 4191   917]
```

```
 [ 4831 23061]]
```


Confusion Matrix




7.4.2 [5.4.2] Applying Logistic Regression with L2 regularization on TFIDF W2V, SET 4

In [256]: *# Please write all the code with proper documentation*

```
all_lr(tfidf_train_vectors,y_train,X_cv,'l2')
```

100%|| 7/7 [00:04<00:00, 1.54it/s]




Cv auc scores with penalty L2

[0.8686668268851596, 0.8946634880571093, 0.8969706221466804, 0.8961376090875506, 0.89598988413]

Maximun Auc value : 0.8969706221466804

Index 2

The optimal value of Lambda = 0.1




Train confusion matrix

```
[[ 6162 1093]
 [ 6475 31160]]
```

Test confusion matrix

```
[[ 4192  916]
 [ 4798 23094]]
```

Confusion Matrix



8 [6] Conclusions

In [257]: *# Please compare all your models using Prettytable library*
from prettytable import PrettyTable

Names of models

Vectorizer = ['Bag of Words', 'Bag of Words', 'TFIDF ', 'TFIDF ', 'AVG W2V', 'AVG W2V', 'T

Param=[0.1, 0.1,0.1,0.1,0.1, 0.1,0.1, 0.1]

auc =[0.89,0.89,0.89,0.89,0.89,0.89,0.89,0.89]

numbering = [1,2,3,4,5,6,7,8]

Initializing prettytable

ptable = PrettyTable()

Adding columns

ptable.add_column("S.NO.", numbering)

ptable.add_column("MODEL", Vectorizer)

ptable.add_column("Hyper Parameter", Param)

```
ptable.add_column("AUC",auc)
```

```
print(ptable)
```

S.NO.	MODEL	Hyper Parameter	AUC
1	Bag of Words	0.1	0.89
2	Bag of Words	0.1	0.89
3	TFIDF	0.1	0.89
4	TFIDF	0.1	0.89
5	AVG W2V	0.1	0.89
6	AVG W2V	0.1	0.89
7	TFIDF W2V	0.1	0.89
8	TFIDF W2V	0.1	0.89

8.0.1 Conclusion

This LR classifier is faster than some of the previous models. TFIDF and BOW both L1 and L2 gave 89% AUC value. The model can be improved by taking more data points, and by taking some other features that may provide us with better insight about the data.