

Amazon-Reviews-on-KNN_Final

March 4, 2019

1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective: Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

2 [1]. Reading Data

2.1 [1.1] Loading the data

The dataset is available in two forms 1. .csv file 2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle
from sklearn.neighbors import KNeighborsClassifier
from sklearn.neighbors import NearestNeighbors

from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score, roc_auc_score

from tqdm import tqdm
import os

In [2]: # using SQLite Table to read data.
con = sqlite3.connect('database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points.
```

```

# you can change the number to any other number based on your computing power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 50000 """)
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 100000 """)

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(-1)
def partition(x):
    if x < 3:
        return 0
    else:
        return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

```

Number of data points in our data (100000, 10)

```

Out[2]:
   Id  ProductId  UserId  ProfileName \
0   1  B001E4KFG0  A3SGXH7AUHU8GW  delmartian
1   2  B00813GRG4  A1D87F6ZCVE5NK  dll pa
2   3  B000LQOCHO  ABXLMWJIXXAIN  Natalia Corres "Natalia Corres"

   HelpfulnessNumerator  HelpfulnessDenominator  Score  Time \
0                      1                      1      1  1303862400
1                      0                      0      0  1346976000
2                      1                      1      1  1219017600

   Summary  Text
0  Good Quality Dog Food  I have bought several of the Vitality canned d...
1  Not as Advertised  Product arrived labeled as Jumbo Salted Peanut...
2  "Delight" says it all  This is a confection that has been around a fe...

```

```

In [3]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)

```

```

In [4]: print(display.shape)
display.head()

```

(80668, 7)

```
Out [4]:
```

	UserId	ProductId	ProfileName	Time	Score	\
0	#oc-R115TNMSPFT9I7	B005ZBZLT4	Breyton	1331510400	2	
1	#oc-R11D9D7SHXIJB9	B005HG9ESG	Louis E. Emory "hoppy"	1342396800	5	
2	#oc-R11DNU2NBKQ23Z	B005ZBZLT4	Kim Cieszykowski	1348531200	1	
3	#oc-R1105J5ZVQE25C	B005HG9ESG	Penguin Chick	1346889600	5	
4	#oc-R12KPBODL2B5ZD	B0070SBEV0	Christopher P. Presta	1348617600	1	

	Text	COUNT(*)
0	Overall its just OK when considering the price...	2
1	My wife has recurring extreme muscle spasms, u...	3
2	This coffee is horrible and unfortunately not ...	2
3	This will be the bottle that you grab from the...	3
4	I didnt like this coffee. Instead of telling y...	2

```
In [5]: display[display['UserId']=='AZY10LLTJ71NX']
```

```
Out [5]:
```

	UserId	ProductId	ProfileName	Time	\
80638	AZY10LLTJ71NX	B001ATMQK2	undertheshrine "undertheshrine"	1296691200	

	Score	Text	COUNT(*)
80638	5	I bought this 6 pack because for the price tha...	5

```
In [6]: display['COUNT(*)'].sum()
```

```
Out [6]: 393063
```

3 [2] Exploratory Data Analysis

3.1 [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [7]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

```
Out [7]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	\
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	

	HelpfulnessDenominator	Score	Time \
0	2	5	1199577600
1	2	5	1199577600
2	2	5	1199577600
3	2	5	1199577600
4	2	5	1199577600

	Summary \
0	LOACKER QUADRATINI VANILLA WAFERS
1	LOACKER QUADRATINI VANILLA WAFERS
2	LOACKER QUADRATINI VANILLA WAFERS
3	LOACKER QUADRATINI VANILLA WAFERS
4	LOACKER QUADRATINI VANILLA WAFERS

	Text
0	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
3	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
4	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8) ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [8]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)
```

```
In [9]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first')
final.shape
```

```
Out[9]: (87775, 10)
```

```
In [10]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

```
Out[10]: 87.775
```

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [11]: display= pd.read_sql_query("""
```

```
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)
```

```
display.head()
```

```
Out[11]:
```

	Id	ProductId	UserId	ProfileName	\
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens	"Jeanne"
1	44737	B001EQ55RW	A2V0I904FH7ABY		Ram

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0		3	1	5	1224892800
1		3	2	4	1212883200

	Summary	\
0	Bought This for My Son at College	
1	Pure cocoa taste with crunchy almonds inside	

	Text
0	My son loves spaghetti so I didn't hesitate or...
1	It was almost a 'love at first bite' - the per...

```
In [12]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [13]: #Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)
```

```
#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()
```

```
(87773, 10)
```

```
Out[13]: 1    73592
0     14181
Name: Score, dtype: int64
```

```
In [14]: ##Sorting data for Time Based Splitting
```

```
time_sorted_data = final.sort_values('Time', axis=0, ascending=True, inplace=False, k
```

```
final = time_sorted_data.take(np.random.permutation(len(final))[:50000])
print(final.shape)
final.head()
```

(50000, 10)

```
Out[14]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	\
83782	91170	B000P3ZJ2Q	A2RUVBPNRYICVB	Deborah1018	0	
91251	99262	B0009ETA76	A3FQ9T2G2GMRX1	Carly	0	
87453	95191	B003UEKFSO	AC299UUMA6MYG	zen	1	
71821	78193	B001BM399U	A3LWIP2I6550TB	X. Nguyen	0	
51348	55785	B001W40ALW	A36G22F9S9DE91	Jackie	1	

	HelpfulnessDenominator	Score	Time	Summary	\
83782	0	1	1337731200	Deborah1018	
91251	0	1	1244419200	Increased milk supply	
87453	1	1	1310860800	zen	
71821	0	1	1320364800	Best baby food out there!	
51348	1	1	1269561600	Best Olives Ever!	

	Text
83782	they where exactly what they said I sent them ...
91251	My doctor recommended that I try Fenugreek to ...
87453	i don't like black tea because they have a bit...
71821	Earth's Best makes the BEST organic baby food ...
51348	My daughter and I are addicted to these olives...

4 [3] Preprocessing

4.1 [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [15]: # printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)
```

```

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)

```

```

they where exactly what they said I sent them to my mom for mothers day and she loved them they
=====
Definetly my favorite coffee sold by Amazon. Just bought 4 lbs. for 8 bucks a pound which is my
=====
This hot sauce is no joke. I've loved spicy food since before I could walk, and can't remember
=====
How can I know an expiry date for this Jams?<br /><br />I couldn't find a date in any bottle.<br />
=====

```

```

In [16]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
sent_150 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)

```

```

they where exactly what they said I sent them to my mom for mothers day and she loved them they

```

```

In [17]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)

```



```

print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)

```

```

they where exactly what they said I sent them to my mom for mothers day and she loved them they
=====
Definetly my favorite coffee sold by Amazon. Just bought 4 lbs. for 8 bucks a pound which is my
=====
This hot sauce is no joke. I've loved spicy food since before I could walk, and can't remember
=====
How can I know an expiry date for this Jams?I couldn't find a date in any bottle.Please let me

```

```

In [18]: # https://stackoverflow.com/a/47091490/4084039
import re

```

```

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase

```

```

In [19]: sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("="*50)

```

```

This hot sauce is no joke. I have loved spicy food since before I could walk, and can not rem
=====

```

```

In [20]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub(r"\S*\d\S*", "", sent_0).strip()
print(sent_0)

```

```

they where exactly what they said I sent them to my mom for mothers day and she loved them they

```

```
In [21]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)
```

This hot sauce is no joke I have loved spicy food since before I could walk and can not remember

```
In [22]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have been removed in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves',
                'you'll', 'you'd', 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
                'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself',
                'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', 'that',
                'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had',
                'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as',
                'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through',
                'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over',
                'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any',
                'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too',
                's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'n',
                've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",
                "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mi',
                "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't",
                'won', "won't", 'wouldn', "wouldn't"])
```

```
In [23]: # Combining all the above students
from tqdm import tqdm
prepr_rev = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
    prepr_rev.append(sentence.strip())
```

100%|| 50000/50000 [00:24<00:00, 2069.51it/s]

```
In [24]: prepr_rev[1500]
```

```
Out[24]: 'hot sauce no joke loved spicy food since could walk not remember ever eating anything'
```

```
In [25]: print(len(prepr_rev))
        final.shape
```

50000

```
Out[25]: (50000, 10)
```

```
In [26]: final ['prepr_rev']= prepr_rev
        final.head(5)
```

```
Out[26]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	\
83782	91170	B000P3ZJ2Q	A2RUVBPNRYICVB	Deborah1018	0	
91251	99262	B0009ETA76	A3FQ9T2G2GMRX1	Carly	0	
87453	95191	B003UEKFS0	AC299UUMA6MYG	zen	1	
71821	78193	B001BM399U	A3LWIP2I6550TB	X. Nguyen	0	
51348	55785	B001W40ALW	A36G22F9S9DE91	Jackie	1	

	HelpfulnessDenominator	Score	Time	Summary	\
83782	0	1	1337731200	Deborah1018	
91251	0	1	1244419200	Increased milk supply	
87453	1	1	1310860800	zen	
71821	0	1	1320364800	Best baby food out there!	
51348	1	1	1269561600	Best Olives Ever!	

	Text	\
83782	they where exactly what they said I sent them ...	
91251	My doctor recommended that I try Fenugreek to ...	
87453	i don't like black tea because they have a bit...	
71821	Earth's Best makes the BEST organic baby food ...	
51348	My daughter and I are addicted to these olives...	

	prepr_rev
83782	exactly said sent mom mothers day loved lasted...
91251	doctor recommended try fenugreek increase milk...
87453	not like black tea bitter taste love tea reall...
71821	earth best makes best organic baby food ever f...
51348	daughter addicted olives blue cheese clearly b...

```
In [27]: # store final table into an SQLite table for future.
        conn = sqlite3.connect('final.sqlite')
        c=conn.cursor()
        conn.text_factory = str
        final.to_sql('Reviews', conn, schema=None, if_exists='replace', index=True, index_label='id')
        conn.close()
        #Loading data
        conn = sqlite3.connect('final.sqlite')
        data=pd.read_sql_query("""select * from Reviews""",conn)
```

5 [4] Featurization

5.1 [4.1] BAG OF WORDS

```
In [28]: from sklearn.model_selection import train_test_split
         #splitting data into Train, C.V and Test
         X_train, X_test, y_train, y_test = train_test_split(final ['prepr_rev'], final['Score']
         X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test_size=0.33)
         print("Train:",X_train.shape,y_train.shape)
         print("CV:",X_cv.shape,y_cv.shape)
         print("Test:",X_test.shape,y_test.shape)
```

```
Train: (22445,) (22445,)
CV: (11055,) (11055,)
Test: (16500,) (16500,)
```

```
In [29]: vectorizer = CountVectorizer(min_df=10, max_features=500)
         vectorizer.fit(X_train)
         #vectorizer.fit(X_train) # fit has to happen only on train data
         # we use the fitted CountVectorizer to convert the text to vector
         X_train_bow = vectorizer.transform(X_train)
         X_cv_bow = vectorizer.transform(X_cv)
         X_test_bow = vectorizer.transform(X_test)
         print("After vectorizations")
         print(X_train_bow.shape, y_train.shape)
         print(X_cv_bow.shape, y_cv.shape)
         print(X_test_bow.shape, y_test.shape)
```

```
After vectorizations
(22445, 500) (22445,)
(11055, 500) (11055,)
(16500, 500) (16500,)
```

5.2 [4.3] TF-IDF

```
In [30]: vect = TfidfVectorizer(ngram_range=(1,2), min_df=10, max_features=500)
         tf_idf_vect = vect.fit(X_train)
         # we use the fitted CountVectorizer to convert the text to vector
         X_train_tfidf = tf_idf_vect.transform(X_train)
         X_cv_tfidf = tf_idf_vect.transform(X_cv)
         X_test_tfidf = tf_idf_vect.transform(X_test)
         print("After vectorizations")
         print(X_train_tfidf.shape, y_train.shape)
         print(X_cv_tfidf.shape, y_cv.shape)
         print(X_test_tfidf.shape, y_test.shape)
```

```

After vectorizations
(22445, 500) (22445,)
(11055, 500) (11055,)
(16500, 500) (16500,)

```

5.3 [4.4] Word2Vec

```

In [31]: # List of sentence in X_train text
sent_of_train=[]
for sent in X_train:
    sent_of_train.append(sent.split())

# List of sentence in X_test text
sent_of_test=[]
for sent in X_test:
    sent_of_test.append(sent.split())

# Train your own Word2Vec model using your own train text corpus
# min_count = 5 considers only words that occurred atleast 5 times
w2v_model=Word2Vec(sent_of_train,min_count=5,size=50, workers=4)
print(w2v_model.wv.most_similar('great'))
print('='*50)
print(w2v_model.wv.most_similar('worst'))

w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_words))

[('good', 0.8064497113227844), ('excellent', 0.7870410680770874), ('awesome', 0.77118951082229)
=====
[('smoothest', 0.8289793729782104), ('horrible', 0.7650426626205444), ('ever', 0.7644084692001)
number of words that occurred minimum 5 times 9150

In [32]: w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_words))
print("sample words ", w2v_words[0:50])

number of words that occurred minimum 5 times 9150
sample words ['bars', 'tasty', 'got', 'low', 'sales', 'price', 'much', 'less', 'overpriced',

```

5.4 [4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

[4.4.1.1] Avg W2v

```

In [33]: i=0
list_of_sentence_cv=[]
for sentence in X_cv:
    list_of_sentence_cv.append(sentence.split())

```

```

In [34]: # average Word2Vec
         # compute average word2vec for each review.
sent_vectors_cv = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_sentence_cv): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need t
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_cv.append(sent_vec)
sent_vectors_cv = np.array(sent_vectors_cv)
print(sent_vectors_cv.shape)
print(sent_vectors_cv[0])

```

100%|| 11055/11055 [00:23<00:00, 472.46it/s]

(11055, 50)

```

[ 0.11423624  0.52322548  0.92737992  0.42657398  0.03149832 -0.49709554
 -0.22744501  0.77123317 -0.8727156   0.49516902  0.65777033 -0.70547067
 -0.31035827  0.61977974  0.3090385   0.31754157  0.11933282  0.53610293
 -0.14034547 -0.36208856 -0.5970275   0.14536198  0.69747322 -0.05873975
  0.4705957   -0.96374629 -0.80560572  0.76205894  0.02416955 -0.26141378
  1.30679555  0.44071176 -0.02415145 -0.39271318  0.37185064  0.41067095
  1.04710724 -0.16281124 -0.99547324 -1.69104051  0.27877323 -0.03820085
  0.34624931  0.13389229  0.0311872   0.67670954 -0.30030563 -0.65574397
 -1.32690578 -0.10916533]

```

```

In [35]: # compute average word2vec for X_test .
test_vectors = [];
for sent in tqdm(sent_of_test):
    sent_vec = np.zeros(50)
    cnt_words = 0;
    for word in sent: #
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    test_vectors.append(sent_vec)

```

```
test_vectors = np.array(test_vectors)
```

```
print(test_vectors.shape)
```

```
print(test_vectors[0])
```

```
100%|| 16500/16500 [00:35<00:00, 463.12it/s]
```

```
(16500, 50)
```

```
[-0.21650994  0.20763936  0.23225507  0.43245629 -0.03147818 -0.34158848
 -0.23489848  0.12679867  0.15458322  0.52804215 -0.32813715 -0.1358679
 -0.02285376 -0.22915929 -0.2054066  0.33587438 -0.28956964 -0.29684678
  0.05607984 -0.19122175  0.05548987  0.22762442  0.54389045 -0.32156981
 -0.21385965 -0.45756952 -0.09783928  0.14551443 -0.48427639  0.15391132
  0.69929648 -0.1097943  -0.14275275 -0.20143936 -0.25195681  0.17130008
  0.31810653 -0.36715548 -0.56052186 -0.52939431  0.56273055 -0.05734205
 -0.01385809 -0.34118139  0.05820298  0.54194045 -0.06114813  0.19944374
 -0.45926711  0.22486217]
```

```
In [36]: # compute average word2vec for X_train .
```

```
train_vectors = [];
```

```
for sent in tqdm(sent_of_train):
```

```
    sent_vec = np.zeros(50)
```

```
    cnt_words = 0;
```

```
    for word in sent: #
```

```
        if word in w2v_words:
```

```
            vec = w2v_model.wv[word]
```

```
            sent_vec += vec
```

```
            cnt_words += 1
```

```
    if cnt_words != 0:
```

```
        sent_vec /= cnt_words
```

```
    train_vectors.append(sent_vec)
```

```
train_vectors = np.array(train_vectors)
```

```
print(train_vectors.shape)
```

```
print(train_vectors[0])
```

```
100%|| 22445/22445 [00:47<00:00, 467.78it/s]
```

```
(22445, 50)
```

```
[ 0.23795041  0.56482501  0.23975698  0.47662384 -0.03241647 -0.42026585
 -0.0920234  0.60887131  0.10361548  0.69549562 -0.42657898 -0.22850905
  0.10854015 -0.43730003 -0.22391569 -0.11131373 -0.21342933 -0.09610408
 -0.10915289 -0.49629217 -0.10829031  0.36527382  0.78678009 -0.25987415
 -0.59941445  0.05826041 -0.38378821  0.2704506  -0.57303713 -0.18668537]
```

```

0.36423239 -0.03373933 -0.12033268 0.11460553 -0.0766927 -0.05885986
0.27646326 -0.3050041 -0.4903128 -0.87627939 1.16383488 0.10803538
-0.10990737 -0.79267004 0.08959021 0.62530343 -0.51399749 0.20063627
-0.29426406 0.25943958]

```

[4.4.1.2] TFIDF weighted W2v

```
In [37]: tf_idf_vect = TfidfVectorizer()
```

```

# final_tf_idf1 is the sparse matrix with row= sentence, col=word and cell_val = tfidf
final_tf_idf1 = tf_idf_vect.fit_transform(X_train)
dictionary = dict(zip(tf_idf_vect.get_feature_names(), list(tf_idf_vect.idf_)))

# tfidf words/col-names
tfidf_feat = tf_idf_vect.get_feature_names()

# compute TFIDF Weighted Word2Vec for X_test .
tfidf_test_vectors = [];
row=0;
for sent in tqdm(sent_of_test):
    sent_vec = np.zeros(50)
    weight_sum =0;
    for word in sent:
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_test_vectors.append(sent_vec)

tfidf_test_vectors = np.array(tfidf_test_vectors)
print(tfidf_test_vectors.shape)
print(tfidf_test_vectors[0])

```

```
100%|| 16500/16500 [06:58<00:00, 39.43it/s]
```

```
(16500, 50)
```

```

[-0.43023417 0.08085112 0.26280191 0.12586552 0.19083609 -0.18514371
-0.28463129 0.13400232 0.16023318 0.52938326 -0.43975282 -0.12440866
-0.11275344 -0.11934448 -0.22762588 0.45164176 -0.39158681 -0.31570446
0.02818469 -0.17517456 0.02748073 0.26977343 0.52292236 -0.34498091

```



```

-0.00175078 -0.44338077  0.02350666  0.06165066 -0.34236525  0.21216029
 0.78613459 -0.15390468 -0.16417371 -0.32861782 -0.2659092  0.10186823
 0.22447934 -0.40924459 -0.62742319 -0.46896159  0.32371111 -0.00329997
-0.03767758 -0.31543755  0.09439429  0.46229863  0.03195479  0.27520308
-0.57627453  0.09213921]

```

```

In [38]: # compute TFIDF Weighted Word2Vec for X_train .
tfidf_train_vectors = [];
row=0;
for sent in tqdm(sent_of_train):
    sent_vec = np.zeros(50)
    weight_sum =0;
    for word in sent:
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_train_vectors.append(sent_vec)

tfidf_train_vectors = np.array(tfidf_train_vectors)
print(tfidf_train_vectors.shape)
print(tfidf_train_vectors[0])

```

100%|| 22445/22445 [07:21<00:00, 50.83it/s]

```

(22445, 50)
[ 0.40917395  0.51609145  0.30560965  0.35250069  0.00734301 -0.39421347
 -0.09543645  0.728009   -0.09724581  0.85466854 -0.43076187 -0.13585101
 -0.06857081 -0.5797596  -0.27492372 -0.41646634 -0.09126814 -0.117861
 -0.0479078  -0.42574795 -0.31589024  0.4247003   0.86006957 -0.2709104
 -0.50717747  0.32116871 -0.42186319  0.35714914 -0.34300058 -0.36286456
  0.35645985 -0.0755837   0.0781704   0.18814762 -0.08119907 -0.23603788
  0.25133437 -0.37718325 -0.22148637 -0.93937279  1.2041414  0.34071557
 -0.01912081 -1.03047928  0.01094786  0.84018511 -0.55510452  0.20887449
 -0.40060992  0.07882321]

```

6 [5] Assignment 3: KNN

Apply Knn(brute force version) on these feature sets

SET 1:Review text, preprocessed one converted into vectors

SET 2:Review text, preprocessed one converted into vectors

- SET 3: Review text, preprocessed one converted into vectors
 - SET 4: Review text, preprocessed one converted into vectors

- SET 5: Review text, preprocessed one converted into vectors

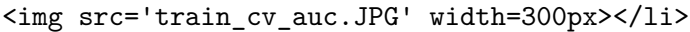
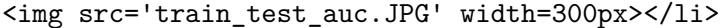
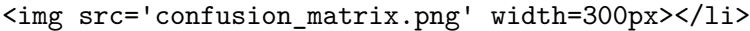
```

count_vect = CountVectorizer(min_df=10, max_features=500)
count_vect.fit(preprocessed_reviews)

```
- SET 6: Review text, preprocessed one converted into vectors

```

tf_idf_vect = TfidfVectorizer(min_df=10, max_features=500)
tf_idf_vect.fit(preprocessed_reviews)

```
- SET 3: Review text, preprocessed one converted into vectors
- SET 4: Review text, preprocessed one converted into vectors
- Find the best hyper parameter which will give the maximum <https://www.appliedaicom>
 - Find the best hyper parameter using k-fold cross validation or simple cross validation data
 - Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this task
- Representation of results
- You need to plot the performance of model both on train data and cross validation data for
 - 
 - Once after you found the best hyper parameter, you need to train your model with it, and find
 - 
 - Along with plotting ROC curve, you need to print the <https://www.appliedaicom>
 - 
- Conclusion
- You need to summarize the results at the end of the notebook, summarize it in the table for

```

    <img src='summary.JPG' width=400px>
</li>
</ul>

```

Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakage, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method `fit_transform()` on your train data, and apply the method `transform()` on cv/test data.
4. For more details please go through this link.

6.1 [5.1] Applying KNN brute force

6.1.1 [5.1.1] Applying KNN brute force on BOW, SET 1

```

In [130]: # https://scikit-learn.org/stable/modules/generated/sklearn.model\_selection.GridSearchCV
          from sklearn.model_selection import GridSearchCV


neigh = KNeighborsClassifier(algorithm='brute')
parameters = {'n_neighbors':[1, 5, 10, 15, 21, 31, 41, 51]}
clf = GridSearchCV(neigh, parameters, cv=3, scoring='roc_auc')
clf.fit(X_train_bow, y_train)

train_auc= clf.cv_results_['mean_train_score']
train_auc_std= clf.cv_results_['std_train_score']
cv_auc = clf.cv_results_['mean_test_score']
cv_auc_std= clf.cv_results_['std_test_score']

plt.plot(K, train_auc, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(K,train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2,color='darkred')

plt.plot(K, cv_auc, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(K,cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2,color='darkred')
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

```



```
In [218]: best_k=10
```

```
In [219]: neigh = KNeighborsClassifier(n_neighbors=best_k)
          neigh.fit(X_train_bow, y_train)
```

```
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of
# not the predicted outputs
```

```
train_fpr, train_tpr, thresholds = roc_curve(y_train, neigh.predict_proba(X_train_bow)[:,1])
test_fpr, test_tpr, thresholds = roc_curve(y_test, neigh.predict_proba(X_test_bow)[:,1])
```

```
plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```

```
print("="*100)
```

```
from sklearn.metrics import confusion_matrix
```

```


print("Train confusion matrix")
print(confusion_matrix(y_train, neigh.predict(X_train_bow)))
print("Test confusion matrix")
print(confusion_matrix(y_test, neigh.predict(X_test_bow)))

```

```

# Variables for table
bow_brute_K = best_k
bow_brute_train= train_auc
bow_brute_test = cv_auc

```



```

=====
Train confusion matrix
[[ 1524  2125]
 [ 1132 17664]]
Test confusion matrix
[[  882  1777]
 [  981 12860]]

```

6.1.2 [5.1.2] Applying KNN brute force on TFIDE, SET 2

```

In [105]: # https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV
from sklearn.model_selection import GridSearchCV

```

```


neigh = KNeighborsClassifier(algorithm='brute')
parameters = {'n_neighbors':[1, 5, 10, 15, 21, 31, 41, 51]}
clf = GridSearchCV(neigh, parameters, cv=3, scoring='roc_auc')
clf.fit(X_train_tfidf, y_train)

train_auc= clf.cv_results_['mean_train_score']
train_auc_std= clf.cv_results_['std_train_score']
cv_auc = clf.cv_results_['mean_test_score']
cv_auc_std= clf.cv_results_['std_test_score']

plt.plot(K, train_auc, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(K,train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2)

plt.plot(K, cv_auc, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(K,cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2,color='darkorange')
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

```



```
In [220]: best_k = 50
```

```
In [221]: neigh = KNeighborsClassifier(n_neighbors=best_k)
neigh.fit(X_train_tfidf, y_train)
```

```
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of
# not the predicted outputs
```

```
train_fpr, train_tpr, thresholds = roc_curve(y_train, neigh.predict_proba(X_train_tfidf)[:,1])
test_fpr, test_tpr, thresholds = roc_curve(y_test, neigh.predict_proba(X_test_tfidf)[:,1])
```


```
plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```

```
print("="*100)
```

```
from sklearn.metrics import confusion_matrix
print("Train confusion matrix")
print(confusion_matrix(y_train, neigh.predict(X_train_tfidf)))
print("Test confusion matrix")
print(confusion_matrix(y_test, neigh.predict(X_test_tfidf)))
```

```
# Variables for table
```

```
tfidf_brute_K = best_k
tfidf_brute_train = train_auc
tfidf_brute_test = cv_auc
```



Train confusion matrix

```
[[ 1 3648]
 [ 0 18796]]
```

Test confusion matrix

```
[[ 1 2658]
 [ 0 13841]]
```

6.1.3 [5.1.3] Applying KNN brute force on AVG W2V, SET 3

```
In [48]: # https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV
from sklearn.model_selection import GridSearchCV

neigh = KNeighborsClassifier(algorithm='brute')
parameters = {'n_neighbors':[1, 5, 10, 15, 21, 31, 41, 51]}
clf = GridSearchCV(neigh, parameters, cv=3, scoring='roc_auc')
clf.fit(train_vectors, y_train)

train_auc= clf.cv_results_['mean_train_score']
train_auc_std= clf.cv_results_['std_train_score']
cv_auc = clf.cv_results_['mean_test_score']
cv_auc_std= clf.cv_results_['std_test_score']
```




```

plt.plot(K, train_auc, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(K,train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2,color='darkblue')

plt.plot(K, cv_auc, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(K,cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2,color='darkorange')
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

```



```
In [222]: best_k = 15
```

```
In [223]: neigh = KNeighborsClassifier(n_neighbors=best_k)
neigh.fit(train_vectors, y_train)
```

```

# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of
# not the predicted outputs

```

```

train_fpr, train_tpr, thresholds = roc_curve(y_train, neigh.predict_proba(train_vectors)[:,1])
test_fpr, test_tpr, thresholds = roc_curve(y_test, neigh.predict_proba(test_vectors)[:,1])

```

```


plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="train AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

print("="*100)

from sklearn.metrics import confusion_matrix
print("Train confusion matrix")
print(confusion_matrix(y_train, neigh.predict(train_vectors)))
print("Test confusion matrix")
print(confusion_matrix(y_test, neigh.predict(test_vectors)))

# Variables for table
Avg_Word2Vec_brute_K = best_k
Avg_Word2Vec_brute_train = train_auc
Avg_word2Vec_brute_test = cv_auc

```



```

Train confusion matrix
[[ 1135  2514]
 [   333 18463]]
Test confusion matrix
[[   706  1953]
 [   304 13537]]

```

6.1.4 [5.1.4] Applying KNN brute force on TFIDF W2V, SET 4

```

In [53]: # https://scikit-learn.org/stable/modules/generated/sklearn.model\_selection.GridSearchCV
         from sklearn.model_selection import GridSearchCV


neigh = KNeighborsClassifier(algorithm='brute')
parameters = {'n_neighbors':[1, 5, 10, 15, 21, 31, 41, 51]}
clf = GridSearchCV(neigh, parameters, cv=3, scoring='roc_auc')
clf.fit(tfidf_train_vectors, y_train)

train_auc= clf.cv_results_['mean_train_score']
train_auc_std= clf.cv_results_['std_train_score']
cv_auc = clf.cv_results_['mean_test_score']
cv_auc_std= clf.cv_results_['std_test_score']

plt.plot(K, train_auc, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(K,train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2)

plt.plot(K, cv_auc, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(K,cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2,color='darkred')
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

```



```
In [224]: best_k=14
```

```
In [225]: neigh = KNeighborsClassifier(n_neighbors=best_k)
neigh.fit(tfidf_train_vectors, y_train)
```

```
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of
# not the predicted outputs
```

```
train_fpr, train_tpr, thresholds = roc_curve(y_train, neigh.predict_proba(tfidf_train_vectors))
test_fpr, test_tpr, thresholds = roc_curve(y_test, neigh.predict_proba(tfidf_test_vectors))
```

```
plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```

```
print("="*100)
```

```
from sklearn.metrics import confusion_matrix
print("Train confusion matrix")
print(confusion_matrix(y_train, neigh.predict(tfidf_train_vectors)))
```

```

print("Test confusion matrix")
print(confusion_matrix(y_test, neigh.predict(tfidf_test_vectors)))


# Variables for table

```

```

TFIDF_Word2Vec_brute_K = best_k
TFIDF_Word2Vec_brute_train = train_auc
TFIDF_word2Vec_brute_test = cv_auc

```



```

=====
Train confusion matrix
[[ 1192  2457]
 [  429 18367]]
Test confusion matrix
[[  702  1957]
 [  390 13451]]

```

6.2 [5.2] Applying KNN kd-tree

6.2.1 [5.2.1] Applying KNN kd-tree on BOW, SET 5

```
In [56]: # Please write all the code with proper documentation
from sklearn.decomposition import TruncatedSVD
svd = TruncatedSVD(n_components=100)
Xtrain_bow = svd.fit_transform(X_train_bow)
Xtest_bow = svd.transform(X_test_bow)


# Please write all the code with proper documentation
# Importing libraries
from sklearn.model_selection import GridSearchCV

neigh = KNeighborsClassifier(algorithm='kd_tree')
parameters = {'n_neighbors':[1, 5, 10, 15, 21, 31, 41, 51]}
clf = GridSearchCV(neigh, parameters, cv=3, scoring='roc_auc')
clf.fit(Xtrain_bow, y_train)

train_auc= clf.cv_results_['mean_train_score']
train_auc_std= clf.cv_results_['std_train_score']
cv_auc = clf.cv_results_['mean_test_score']
cv_auc_std= clf.cv_results_['std_test_score']

plt.plot(K, train_auc, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(K,train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2)

plt.plot(K, cv_auc, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(K,cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2,color='darkred')
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```



```
In [226]: best_k=41
```

```
In [227]: neigh = KNeighborsClassifier(n_neighbors=best_k)
neigh.fit(Xtrain_bow, y_train)
```

```
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of
# not the predicted outputs
```

```
train_fpr, train_tpr, thresholds = roc_curve(y_train, neigh.predict_proba(Xtrain_bow)[:,1])
test_fpr, test_tpr, thresholds = roc_curve(y_test, neigh.predict_proba(Xtest_bow)[:,1])
```


```
plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```

```
print("="*100)
```

```
from sklearn.metrics import confusion_matrix
print("Train confusion matrix")
print(confusion_matrix(y_train, neigh.predict(Xtrain_bow)))
```

```
print("Test confusion matrix")
print(confusion_matrix(y_test, neigh.predict(Xtest_bow)))
```

```
# Variables for table
bow_kdTree_K = best_k
bow_kdTree_train = train_auc
bow_kdTree_test = cv_auc
```



```
=====
Train confusion matrix
[[ 793 2856]
 [ 487 18309]]
Test confusion matrix
[[ 532 2127]
 [ 393 13448]]
```

6.2.2 [5.2.2] Applying KNN kd-tree on TFIDE, SET 6

```
In [60]: # Please write all the code with proper documentation
svd = TruncatedSVD(n_components=100)
Xtrain_tfidf = svd.fit_transform(X_train_bow)
```



```

Xtest_tfidf = svd.transform(X_test_bow)

# Please write all the code with proper documentation
# Importing libraries
from sklearn.model_selection import GridSearchCV


neigh = KNeighborsClassifier(algorithm='kd_tree')
parameters = {'n_neighbors':[1, 5, 10, 15, 21, 31, 41, 51]}
clf = GridSearchCV(neigh, parameters, cv=3, scoring='roc_auc')
clf.fit(Xtrain_tfidf, y_train)

train_auc= clf.cv_results_['mean_train_score']
train_auc_std= clf.cv_results_['std_train_score']
cv_auc = clf.cv_results_['mean_test_score']
cv_auc_std= clf.cv_results_['std_test_score']

plt.plot(K, train_auc, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(K,train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2)

plt.plot(K, cv_auc, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(K,cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2,color='darkred')
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

```



```
In [228]: best_k=31
```

```
In [229]: neigh = KNeighborsClassifier(n_neighbors=best_k)
          neigh.fit(Xtrain_tfidf, y_train)
```

```
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of
# not the predicted outputs
```

```
train_fpr, train_tpr, thresholds = roc_curve(y_train, neigh.predict_proba(Xtrain_tfidf)[:,1])
test_fpr, test_tpr, thresholds = roc_curve(y_test, neigh.predict_proba(Xtest_tfidf)[:,1])
```

```
plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```

```
print("="*100)
```


```
from sklearn.metrics import confusion_matrix
```

```

print("Train confusion matrix")
print(confusion_matrix(y_train, neigh.predict(Xtrain_tfidf)))
print("Test confusion matrix")
print(confusion_matrix(y_test, neigh.predict(Xtest_tfidf)))

# Variables for table
tfidf_kdTree_K = best_k
tfidf_kdTree_train = train_auc
tfidf_kdTree_test = cv_auc

```



```

=====
Train confusion matrix
[[ 850 2799]
 [ 537 18259]]
Test confusion matrix
[[ 567 2092]
 [ 450 13391]]

```

6.2.3 [5.2.3] Applying KNN kd-tree on AVG W2V, SET 7

```
In [63]: from sklearn.model_selection import GridSearchCV
```

```


neigh = KNeighborsClassifier(algorithm='kd_tree')
parameters = {'n_neighbors':[1, 5, 10, 15, 21, 31, 41, 51]}
clf = GridSearchCV(neigh, parameters, cv=3, scoring='roc_auc')
clf.fit(train_vectors, y_train)

train_auc= clf.cv_results_['mean_train_score']
train_auc_std= clf.cv_results_['std_train_score']
cv_auc = clf.cv_results_['mean_test_score']
cv_auc_std= clf.cv_results_['std_test_score']

plt.plot(K, train_auc, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(K,train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2)

plt.plot(K, cv_auc, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(K,cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2,color='darkorange')
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

```



In [230]: best_k=15

```

In [231]: neigh = KNeighborsClassifier(n_neighbors=best_k)
          neigh.fit(train_vectors, y_train)

          # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates o
          # not the predicted outputs

          train_fpr, train_tpr, thresholds = roc_curve(y_train, neigh.predict_proba(train_vecto
          test_fpr, test_tpr, thresholds = roc_curve(y_test, neigh.predict_proba(test_vectors)


          plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
          plt.plot(test_fpr, test_tpr, label="train AUC =" + str(auc(test_fpr, test_tpr)))
          plt.legend()
          plt.xlabel("K: hyperparameter")
          plt.ylabel("AUC")
          plt.title("ERROR PLOTS")
          plt.show()

          print("="*100)

          from sklearn.metrics import confusion_matrix
          print("Train confusion matrix")
          print(confusion_matrix(y_train, neigh.predict(train_vectors)))
          print("Test confusion matrix")
          print(confusion_matrix(y_test, neigh.predict(test_vectors)))

          # Variables for table
          Avg_Word2Vec_kdTree_K = best_k
          Avg_Word2Vec_kdTree_train = train_auc
          Avg_Word2Vec_kdTree_test = cv_auc

```



=====

Train confusion matrix

```
[[ 1135  2514]
 [   333 18463]]
```

Test confusion matrix

```
[[   706   1953]
 [   304 13537]]
```

6.2.4 [5.2.4] Applying KNN kd-tree on TFIDF W2V, SET 8

```
In [67]: neigh = KNeighborsClassifier(algorithm='kd_tree')
         parameters = {'n_neighbors':[1, 5, 10, 15, 21, 31, 41, 51]}
         clf = GridSearchCV(neigh, parameters, cv=3, scoring='roc_auc')
         clf.fit(tfidf_train_vectors, y_train)

         train_auc= clf.cv_results_['mean_train_score']
         train_auc_std= clf.cv_results_['std_train_score']
         cv_auc = clf.cv_results_['mean_test_score']
         cv_auc_std= clf.cv_results_['std_test_score']


         plt.plot(K, train_auc, label='Train AUC')
```

```

# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(K,train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2)

plt.plot(K, cv_auc, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(K,cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2,color='darkred')
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

```



```
In [232]: best_k=15
```

```
In [233]: neigh = KNeighborsClassifier(n_neighbors=best_k)
neigh.fit(tfidf_train_vectors, y_train)
```

```

# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of
# not the predicted outputs

```

```

train_fpr, train_tpr, thresholds = roc_curve(y_train, neigh.predict_proba(tfidf_train_vectors)[:,1])
test_fpr, test_tpr, thresholds = roc_curve(y_test, neigh.predict_proba(tfidf_test_vectors)[:,1])

```

```


plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="train AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()

print("="*100)

from sklearn.metrics import confusion_matrix
print("Train confusion matrix")
print(confusion_matrix(y_train, neigh.predict(tfidf_train_vectors)))
print("Test confusion matrix")
print(confusion_matrix(y_test, neigh.predict(tfidf_test_vectors)))

# Variables for table
TFIDF_Word2Vec_kdTree_K = best_k
TFIDF_Word2Vec_kdTree_train = train_auc
TFIDF_Word2Vec_kdTree_test = cv_auc

```



=====

Train confusion matrix


```
[[ 962 2687]
 [ 276 18520]]
Test confusion matrix
[[ 588 2071]
 [ 269 13572]]
```

7 [6] Conclusions

In [235]: `from prettytable import PrettyTable`

```
# Initializing prettytable
```

```
ptable = PrettyTable()
```

```
names = ["brute for BoW", "kdTree for BoW", "brute for TFIDF", "kdTree' for TFIDF", \
        "brute for TFIDF-Word2Vec", "kdTree' for TFIDF-Word2Vec"]
```

```
k = ([bow_brute_K , bow_kdTree_K , tfidf_brute_K , tfidf_kdTree_K , Avg_Word2Vec_brute_K , \
      TFIDF_Word2Vec_brute_K , TFIDF_Word2Vec_kdTree_K ])
```

```
train = ([bow_brute_train, bow_kdTree_train, tfidf_brute_train, tfidf_kdTree_train, \
          Avg_Word2Vec_brute_train, Avg_Word2Vec_kdTree_train, TFIDF_Word2Vec_brute_train, \
          TFIDF_Word2Vec_kdTree_train])
```

```
test = ([bow_brute_test, bow_kdTree_test , tfidf_brute_test , tfidf_kdTree_test , \
         Avg_word2Vec_brute_test, Avg_Word2Vec_kdTree_test , TFIDF_word2Vec_brute_test, \
         TFIDF_Word2Vec_kdTree_test ])
```

```
ptable.add_column("MODEL",names)
ptable.add_column("K: hyperparameter",k)
ptable.add_column("Train",train)
ptable.add_column("Test",test)
```

```
# Printing the Table
```

```
print(ptable)
```

MODEL	K: hyperparameter	Train
brute for BoW	10	[0.99904082 0.87564734 0.83360146 0.81453208 0.78554394 0.78689953]
kdTree for BoW	41	[0.99904082 0.87564734 0.83360146 0.81453208 0.78554394 0.78689953]
brute for TFIDF	50	[0.99904082 0.87564734 0.83360146 0.81453208 0.78554394 0.78689953]
kdTree' for TFIDF	31	[0.99904082 0.87564734 0.83360146 0.81453208 0.78554394 0.78689953]

					0.78554394 0.78689953
	brute for Avg-Word2Vec		15		[0.99904082 0.87564734 0.83360146 0.81453208
					0.78554394 0.78689953
	kdTree for Avg-Word2Vec		15		[0.99904082 0.87564734 0.83360146 0.81453208
					0.78554394 0.78689953
	brute for TFIDF-Word2Vec		14		[0.99904082 0.87564734 0.83360146 0.81453208
					0.78554394 0.78689953
	kdTree' for TFIDF-Word2Vec		15		[0.99904082 0.87564734 0.83360146 0.81453208
					0.78554394 0.78689953
+-----+-----+-----+-----+-----+-----+					

In []: